

Shiv Nadar University Chennai

Degree & Branch	B.Tech Artificial Intelligence & Data Science	Semester V
Subject Code & Name	CS3807 – Deep Learning Laboratory	AY: 2026–27

Experiment 1 Implementation of a Single Layer Perceptron for Binary Classification

1. Objective

The objective of this experiment is to understand the concept of an artificial neuron and implement a Single Layer Perceptron from scratch for binary classification. Students will learn the perceptron learning algorithm, understand the importance of activation functions, visualize the learning process, and evaluate the classifier using a real-world dataset.

2. Tasks to be Performed

1. Download and understand the dataset.
2. Perform Exploratory Data Analysis (EDA).
3. Preprocess the dataset.
4. Implement a Single Layer Perceptron from scratch.
5. Train the model using the perceptron learning rule.
6. Evaluate the classifier using standard performance metrics.
7. Analyze the effect of different learning rates.
8. Compare the implementation with Scikit-learn's Perceptron (optional).

3. Background Theory

An Artificial Neuron is the fundamental building block of a neural network. It receives one or more input features, multiplies them by corresponding weights, adds a bias term, and applies an activation function to determine the output.

The Single Layer Perceptron is the earliest supervised learning algorithm capable of solving linearly separable binary classification problems.

Mathematical Formulation

Weighted Sum

$$z = \mathbf{w}^T \mathbf{x} + b$$

where

- $\mathbf{x}$: Input vector
- $\mathbf{w}$: Weight vector
- b : Bias
- z : Linear combination

Prediction

$$\hat{y} = f(z)$$

Weight Update Rule

$$\mathbf{w}_{new} = \mathbf{w}_{old} + \eta(y - \hat{y})\mathbf{x}$$

Bias Update

$$b_{new} = b_{old} + \eta(y - \hat{y})$$

where η denotes the learning rate.

4. Activation Functions

An activation function determines the output of a neuron based on the weighted sum of its inputs. It introduces non-linearity, enabling neural networks to learn complex patterns. Although the perceptron uses only the Step Activation Function, modern deep learning models employ differentiable activation functions for efficient optimization through backpropagation.

Step Activation Function

$$f(z) = \begin{cases} 1, & z \geq 0 \\ 0, & z < 0 \end{cases}$$

The Step Activation Function produces a binary output and is therefore suitable only for linearly separable binary classification problems.

Activation	Output Range	Typical Applications
Step	$\{0, 1\}$	Classical Perceptron
Sigmoid	$(0, 1)$	Binary Classification Output Layer
Tanh	$(-1, 1)$	Hidden Layers (Traditional Networks)
ReLU	$[0, \infty)$	Hidden Layers in CNNs and MLPs
Leaky ReLU	$(-\infty, \infty)$	Avoiding Dying ReLU
Softmax	$(0, 1)$	Multi-class Classification Output Layer

Note: This experiment uses only the Step Activation Function. The remaining activation functions will be studied in subsequent experiments involving Multilayer Perceptrons and Deep Neural Networks.

5. Dataset Description

Dataset: Banknote Authentication Dataset

Source: UCI Machine Learning Repository

Download Link:

<https://archive.ics.uci.edu/dataset/267/banknote+authentication>

Instances	1372
Features	4 Numerical Features
Classes	2
Missing Values	None
Task	Binary Classification

Feature Description

- Variance
- Skewness
- Curtosis
- Entropy

Target Class

- 0 – Authentic Banknote
- 1 – Forged Banknote

6. Experimental Procedure

Task 1: Dataset Exploration

- Load the dataset.
- Display the first five samples.
- Determine dataset dimensions.
- Identify missing values.
- Display descriptive statistics.

Expected Output

Dataset summary and statistical information.

Observations

The dataset has 1372 rows and 5 columns (4 features + 1 target). There are no missing values anywhere in the data. The mean of variance is around 0.43, skewness is around 1.92, kurtosis is around 1.40 and entropy is around -1.19 . The target column has slightly more 0s (authentic notes) than 1s (forged notes), so the data is fairly balanced.

Task 2: Exploratory Data Analysis

Generate

- Histograms
- Correlation Heatmap
- Scatter Plot
- Boxplots

Task 3: Data Preprocessing

- Normalize all numerical features.
- Split the dataset into Training (80%) and Testing (20%).

Task 4: Perceptron Implementation

Implement from scratch

- Weight Initialization
- Bias Initialization
- Step Activation Function
- Forward Propagation
- Perceptron Learning Rule

Task 5: Model Training

Display

- Epoch Number
- Misclassified Samples
- Updated Weights
- Updated Bias

Task 6: Model Evaluation

Compute

- Accuracy
- Precision
- Recall
- F1-score
- Confusion Matrix

7. Source Code

The complete source code for this experiment, including the dataset loading, EDA, preprocessing, implementation of perceptron, training, evaluation, and all plot's code, is available k `DL_Lab_1.ipynb` on GitHub.

GitHub Repository Link:

<https://github.com/Siddharth-G06/Deep-Learning-Lab/tree/main/Lab%201%20-%20Single%20Layer%20Perceptron>

The notebook can be run end-to-end using the steps mentioned in the README of the repository (install dependencies from `requirements.txt`, then run `jupyter notebook DL_Lab_1.ipynb`).

8. Experimental Results

Dataset Exploration

The dataset had 1372 samples and 4 features, with no missing values anywhere. The class split was fairly balanced – 762 authentic notes and 610 forged notes.

Training Results (Learning Rate = 0.01)

The perceptron was trained for 100 epochs. The number of misclassified samples dropped sharply from 171 in epoch 1 to under 30 by epoch 10, but the model never fully converged to zero errors – it kept oscillating between roughly 16 and 30 misclassified samples for the remaining epochs. Training ended at epoch 100 with:

- Final Weights: $[-0.2126, -0.2234, -0.2247, -0.0067]$
- Final Bias: 0.29
- Misclassified samples in final epoch: 22 (out of 1097 training samples)

Test Set Evaluation

The trained model was evaluated on the 275 held-out test samples, giving the following results:

Metric	Value
Accuracy	0.9891
Precision	0.9917
Recall	0.9836
F1-Score	0.9877

Out of 275 test samples, the model correctly classified 152 authentic notes and 120 forged notes, with only 1 false positive and 2 false negatives.

Learning Rate Comparison

The experiment was repeated with learning rates 0.001 and 0.1, in addition to 0.01. All three learning rates produced identical error curves and identical final test accuracy (0.9891). This happens because the weights were initialized to zero, so the learning rate only scales the size of the weight updates and never changes the sign of the prediction – meaning the same points get classified the same way at every epoch, no matter which learning rate is used.

Learning Rate	Accuracy	Precision	Recall	F1-Score
0.001	0.9891	0.9917	0.9836	0.9877
0.01	0.9891	0.9917	0.9836	0.9877
0.1	0.9891	0.9917	0.9836	0.9877

9. Generated Plots with Interpretations

Plot	Purpose	Expected Inference
Feature Histograms	Distribution Analysis	Identify skewness and outliers.
Correlation Heatmap	Feature Relationship	Observe highly correlated features.
Scatter Plot	Class Distribution	Determine whether classes are approximately linearly separable.
Training Error vs Epoch	Learning Behaviour	Verify convergence of the perceptron.
Weight Evolution	Parameter Analysis	Observe stabilization of learned weights.
Bias Evolution	Parameter Analysis	Understand the role of the bias term.
Confusion Matrix	Performance Evaluation	Interpret TP, FP, TN and FN.
Learning Rate Comparison	Hyperparameter Study	Compare convergence behaviour for different learning rates.
Decision Boundary (Optional)	Visualization	Illustrate the separating hyperplane using two selected features.

Students should provide a brief interpretation (2–3 sentences) for every generated plot.

9.1 Feature Histograms




Figure 1: Histograms of all four features and the target class.

Interpretation: Variance and skewness are more or less spread out evenly, kind of like a normal curve. kurtosis and entropy are not like that — they lean more to one side, entropy especially has a lot of values crowded near the negative side. The target class graph shows we have a bit more authentic notes than forged ones, but not by a huge margin.

9.2 Correlation Heatmap




Figure 2: Correlation heatmap of all features.

Interpretation: Variance has the strongest correlation with the target class, around -0.72 , so this feature alone tells us a lot about whether a note is real or fake. Skewness and kurtosis are strongly negatively correlated with each other (around -0.79), meaning when one goes up the other usually goes down. None of the features are copies of each other, so all four are useful to keep.

9.3 Scatter Plot




Figure 3: Pairwise scatter plot coloured by class.

Interpretation: When we look at variance and skewness together, the blue points (authentic) and orange points (forged) form two separate groups, with only a little overlap in the middle. This tells us the classes are almost linearly separable, which is a good sign for a perceptron since it can only draw straight line boundaries. This overlap is also probably why our model didn't get 100% accuracy.

9.4 Box Plots




Figure 4: Box plots of the four features.

Interpretation: Skewness has the widest range and biggest spread out of all the features. kurtosis has quite a few outlier points sitting above the upper whisker. Since the features have very different ranges before scaling, this is exactly why we needed to normalize everything before training the perceptron.

9.5 Training Error vs Epoch




Figure 5: Misclassified samples per epoch, learning rate = 0.01.

Interpretation: The error drops really fast in the first 10 epochs, from 171 wrong predictions down to under 30. After that point, it doesn't go to zero — it just keeps bouncing between about 16 and 30 errors till epoch 100. This happens because a few points from both classes overlap each other, so the perceptron keeps correcting the same handful of points again and again without ever fully settling down.

9.6 Weight Evolution




Figure 6: Weight values over 100 epochs.

Interpretation: All four weights start at 0 (as we set them) and move down into negative values pretty quickly within the first 15 epochs. After that, they mostly flatten out and just move a little here and there instead of changing a lot. This basically means the model found its rough decision boundary early on and spent the rest of the epochs just adjusting it slightly.

9.7 Bias Evolution




Figure 7: Bias value over 100 epochs.

Interpretation: Unlike the weights, the bias keeps climbing up slowly and steadily throughout training, going from 0 to almost 0.29 by the end, like a staircase. It never really flattens out completely because misclassifications keep happening every epoch, even if the count is small, and each one nudges the bias up a tiny bit more.

9.8 Confusion Matrix




Figure 8: Confusion matrix on test data.

Interpretation: Out of 275 test samples, our model got 152 authentic notes and 120 forged notes correct, and only messed up 3 samples total (1 false positive, 2 false negatives). This means the model works really well on data it has never seen before, not just on the training data.

9.9 Learning Rate Comparison




Figure 9: Error curves for learning rates 0.001, 0.01, 0.1.

Interpretation: Looking at the graph carefully, only the green line ($LR = 0.1$) is visible even though all three learning rates are plotted — this is because all three curves are exactly overlapping each other, not just similar. This happens because our weights start at zero, and the perceptron update rule only scales the weight and bias values by the learning rate — it doesn't change whether a point gets classified correctly or not, since that only depends on the sign of the output, not its size. So changing the learning rate only changes how big the weight numbers get, not the actual predictions made at each epoch. That's exactly why all three learning rates ended up with the identical error curve and the identical final test accuracy of 98.91 percent.

10. Performance Tables

Training Summary

Table 2: Model Training Parameters and Performance

Parameter	Value
Dataset Size	1372 samples
Train/Test Split	1097 / 275 (80/20, stratified)
Learning Rate	0.01
Epochs	100 (did not fully converge, kept oscillating)
Final Weights	[-0.2126, -0.2234, -0.2247, -0.0067]
Final Bias	0.29
Accuracy	0.9891
Precision	0.9917
Recall	0.9836
F1-score	0.9877

Epoch-wise Learning

Table 3: Epoch-wise Learning Progress

Epoch	Errors	Weight 1	Weight 2	Bias
1	171	-0.0679	-0.0462	0.07
25	31	-0.1319	-0.1506	0.20
50	20	-0.1628	-0.1898	0.23
75	18	-0.1871	-0.2114	0.26
100	22	-0.2126	-0.2234	0.29

11. Additional Tasks

1. Compare the Step and Sigmoid activation functions. Discuss why the Step Activation Function is not used in modern deep learning models.

Answer : Step function gives 0 if $f(x) \leq 0$ and 1 when $f(x) > 0$, and it has slope 0 everywhere except at 0 hence we cannot use it with gradient descent, Sigmoid on the other hand gives values between 0 and 1 and has smooth slopes everywhere, so it works with backpropagation, thats why modern deep learning models use sigmoid and not step.

2. Compare the implementation with Scikit-learn's Perceptron.

Answer : Both the ideas are same , which is updating the weights, but the scikit version is more optimized but our version lets us see the error, weight and bias values epoch by epoch.

3. Repeat the experiment using learning rates 0.001, 0.01 and 0.1.

Answer :All three learning rates (0.001, 0.01, 0.1) produced the exact same error curve and the same final test accuracy of 98.91. This is because the perceptron was initialized with zero weights, and the learning rate only scales the size of the weight updates it doesn't affect the sign of the prediction, which is what decides right or wrong. So the three models made identical classification decisions at every epoch, just with different-sized weight numbers underneath. This would not stay true if the weights were initialized to non-zero random values instead of zero. .

4. Explain why a Single Layer Perceptron cannot solve the XOR problem.

Answer : In XOR, the points (0,1) and (1,0) give output 1, and (0,0) and (1,1) give output 0s since these points are placed diagonally opposite, it is not possible for a straight line to seperate them, but a perceptron can only draw a straight boundry, hence a single layer perceptron cannot solve XOR problem and we need multiple layers.

5. Study the effect of feature normalization on model convergence.

Answer : Before normalizing, features like entropy (range -8.5 to 2.4) and skewness (range -14 to 13) were on totally different scales, which would mess up the weight updates unfairly. After Min-Max scaling everything to 0-1, all features got treated equally, and that's a big reason the model reached a low-error zone within just 10-15 epochs..

12. Discussion

The perceptron worked really well on this dataset, getting about 98.9% accuracy on data it never saw during training. But it's worth pointing out that it never actually reached 0 errors during training — the error count kept jumping between 16 and 30 samples out of 1097, even after 100 epochs. This is because the two classes overlap a little bit (seen clearly in the scatter plot), so the data is not perfectly linearly separable, just close to it. The learning rate comparison also showed something interesting — even though the three rates behaved very differently while training, all of them landed on the same final test performance, meaning the learning rate mostly affects the training journey, not the final destination, at least for this dataset.

13. Conclusion

We successfully built a Single Layer Perceptron from scratch, without using any built-in ML classifier, and used it to tell apart real and forged banknotes using 4 simple features. After normalizing

the data and training for 100 epochs at a learning rate of 0.01, the model reached 98.91% accuracy, 99.17% precision, 98.36% recall and an F1-score of 98.77% on the test set. The model didn't fully converge because the two classes overlap slightly, but the results are still very strong for a plain single-layer model. This experiment gave a clear, hands-on understanding of how something as basic as the perceptron learning rule can still solve real classification problems fairly accurately.

12. Report Contents

The report should include:

- Objective
- Background Theory
- Activation Functions
- Dataset Description
- Experimental Procedure
- Source Code
- Experimental Results
- Generated Plots with Interpretation
- Performance Tables
- Discussion
- Conclusion
- References

14. References

1. F. Rosenblatt, "The Perceptron," *Psychological Review*, 1958.
2. I. Goodfellow, Y. Bengio and A. Courville, *Deep Learning*, MIT Press, 2016.
3. C. M. Bishop, *Pattern Recognition and Machine Learning*, Springer, 2006.
4. S. Haykin, *Neural Networks and Learning Machines*, Pearson, 2009.
5. UCI Machine Learning Repository – Banknote Authentication Dataset.
6. Scikit-learn Documentation: Perceptron.